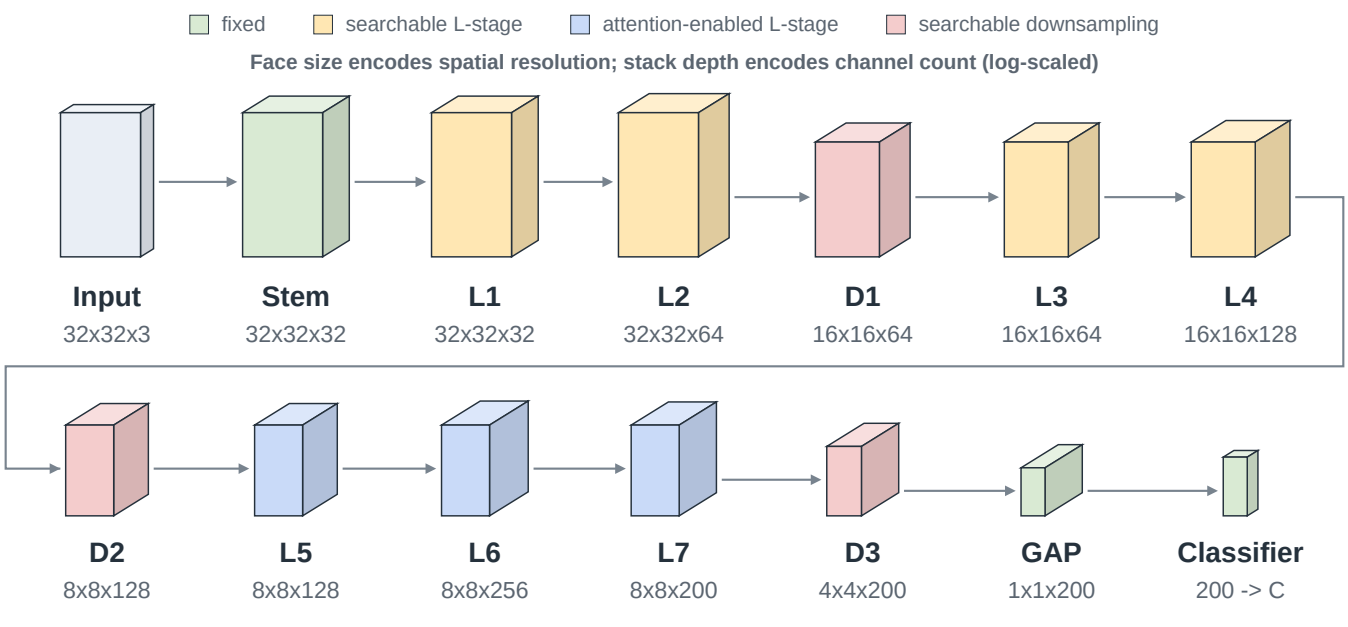
fixed
searchable L-stage
attention-enabled L-stage
searchable downsampling
Face size encodes spatial resolution; stack depth encodes channel count (log-scaled)
Input
32x32x3
Stem
32x32x32
L1
32x32x32
L2
32x32x64
D1
16x16x64
L3
16x16x64
L4
16x16x128
D2
8x8x128
L5
8x8x128
L6
8x8x256
L7
8x8x200
D3
4x4x200
GAP
1x1x200
Classifier
200 -> C